

FIRST-PARTY BUILD REPORT

AGENTIC FLEET CONTROL PLANE

Eight AI employees. One accountable operating system.

A public field guide from Roman Bediner's Fractional COO work

PUBLIC FIELD GUIDE

01 THE CONTROL PLANE

02 THE EMPLOYEES


03 THE ROUTER

04 THE RELEASE LOOP

romanbediner.com/resources/agentic-ai-employees

The stack, end to end

The repo is the body. Hivemind holds shared memory and operating behavior. Each run retrieves the relevant context, acts through clients, leaves a heartbeat, and carries work to surfaces the organization can verify.



Eight named agents. Clear jobs. One human authority boundary.

The names are public. The private channels, records, and operating instructions are not. Each role has a lane, a memory, and a controlled route to action

HUMAN OWNER > DIRECTOR OF FLEET ORCHESTRATION & ENGINEERING

direction, protected decisions, and accountable operation

STAFF ENGINEER

independent review gate

PROJECT MANAGER

work and commitments

CHIEF OF STAFF

knowledge and briefs

REVOPS ENGINEER

revenue intelligence

IMPROVEMENT ENGINEER

learning and backlog

AI CORRESPONDENT

editorial intelligence

SOLUTION ARCHITECT

member proposals

No single agent writes, approves, and ships its own change.

The fleet does not use one model for every job

Model choice becomes an operational policy. Routine work gets speed; judgment and production gates get deeper reasoning. The authority of each job stays visible.

FAST

classification, extraction, routine reports

low cost / low latency

CAPABLE

coordination, synthesis, code changes

balanced reasoning

DEEP

review, recovery, production judgment

reserved for risk

THE RELIABILITY LOOP

Report and recover

HEARTBEAT

did it run?



DIGEST

what happened?




RECOVER

what now?

Reporting is not a dashboard bolted on afterward. It closes the distance between autonomous action and accountable operations.

From engineering queue to verified production

No single agent writes, approves, and ships its own change. The system separates initiative from authority and makes failure recoverable.



The operating consequence: autonomous work can be fast without being opaque.

Build autonomy that stays accountable.

A portable operating model for teams that want more autonomy without losing evidence, ownership, or recovery.

AI is available to everyone.

An operating system for it is not.

A MINIMAL RECIPE

- 1 Define lanes**
Give each agent a job, memory, manager, and authority.
- 2 Make loops visible**
Design work, engineering, and reliability as closed loops.
- 3 Separate initiative / authority**
Let agents prepare; reserve high-downside decisions.
- 4 Leave evidence**
Leave heartbeats, review records, and deployment evidence.
- 5 Expand carefully**
Expand only when reporting and recovery work.